

# Clasificación de imágenes usando Transfer Learning con Data customizada

Leonardo Leon V.

# Outline

- Estado del arte
- Teoría
- Arquitecturas
- Transfer Learning
- Notebook

# Deep Learning in One Slide

- **What is it:**

Extract useful patterns from data.

- **How:**

Neural network + optimization

- **How (Practical):**

Python + TensorFlow & friends

- **Hard Part:**

Good Questions + Good Data

- **Why now:**

Data, hardware, community, tools, investment

- **Where do we stand?**

Most big questions of intelligence have not been answered nor properly formulated

## Exciting progress:

- Face recognition
- Image classification
- Speech recognition
- Text-to-speech generation
- Handwriting transcription
- Machine translation
- Medical diagnosis
- Cars: drivable area, lane keeping
- Digital assistants
- Ads, search, social recommendations
- Game playing with deep RL

# History of Deep Learning Ideas and Milestones\*




## Perspective:

- Universe created  
13.8 billion years ago
- Earth created  
4.54 billion years ago
- Modern humans  
300,000 years ago
- Civilization  
12,000 years ago
- Written record  
5,000 years ago

- 1943: Neural networks
- 1957: Perceptron
- 1974-86: Backpropagation, RBM, RNN
- 1989-98: CNN, MNIST, LSTM, Bidirectional RNN
- 2006: “Deep Learning”, DBN
- 2009: ImageNet
- 2012: AlexNet, Dropout
- 2014: GANs
- 2014: DeepFace
- 2016: AlphaGo
- 2017: AlphaZero, Capsule Networks
- 2018: BERT

\* Dates are for perspective and not as definitive historical record of invention or credit

# Deep Learning Frameworks





## Factors to consider:

- Learning curve
- Speed of development
- Size and passion of community
- Number of papers implemented in framework
- Likelihood of long-term growth and stability
- Ecosystem of tooling

1.  TensorFlow
2.  Keras
3.  PyTorch
4.  Caffe
5.  theano
6.  Apache mxnet™
7.  Microsoft CNTK
8.  DL4J
9.  Caffe2
10.  Chainer
11.  fast.ai

# Why Deep Learning? **Scalable** Machine Learning




**Interpretability of Deep Learning:**

<https://www.youtube.com/watch?v=93Xv8vJ2acI>


# Learning of object parts

Examples of learned object parts from object categories




Faces





Cars





Elephants




Chairs

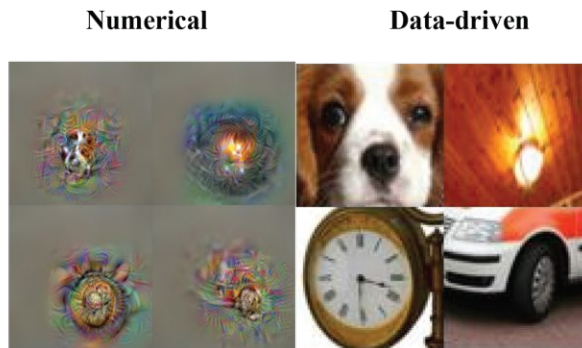




**Conv 1: Edge+Blob**



**Conv 3: Texture**



**Conv 5: Object Parts**



**Fc8: Object Classes**


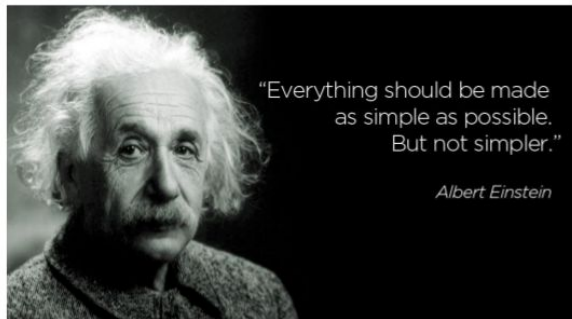


# The Challenge of Deep Learning: Efficient Teaching + Efficient Learning

- Humans can learn from very few examples
- Machines (in most cases) need thousands/millions of examples



# First Steps: Start Simple



1

```
# import tensorflow and keras (tf.keras not "vanilla" Keras)
import tensorflow as tf
from tensorflow import keras
```

2

```
# get data
(train_images, train_labels), (test_images, test_labels) = \
keras.datasets.mnist.load_data()
```

3

```
# setup model
model = keras.Sequential([
    keras.layers.Flatten(input_shape=(28, 28)),
    keras.layers.Dense(128, activation=tf.nn.relu),
    keras.layers.Dense(10, activation=tf.nn.softmax)
])

model.compile(optimizer=tf.train.AdamOptimizer(),
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
```

4

```
# train model
model.fit(train_images, train_labels, epochs=5)
```



5

```
# evaluate
test_loss, test_acc = model.evaluate(test_images, test_labels)
print('test accuracy:', test_acc)
```


6

```
# make predictions
predictions = model.predict(test_images)
```

# Convolutional Layer

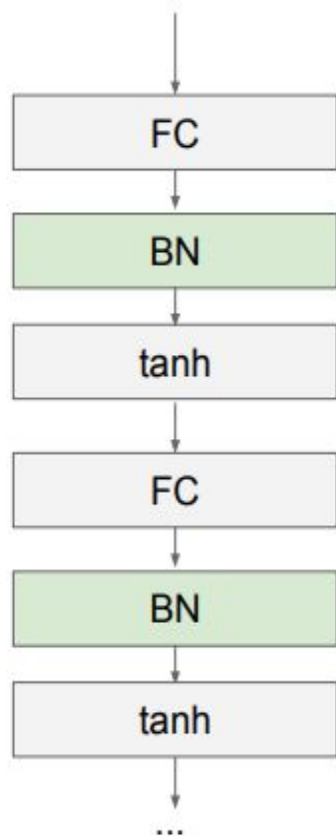


# Pooling Layer



# Batch Normalization


[Ioffe and Szegedy, 2015]



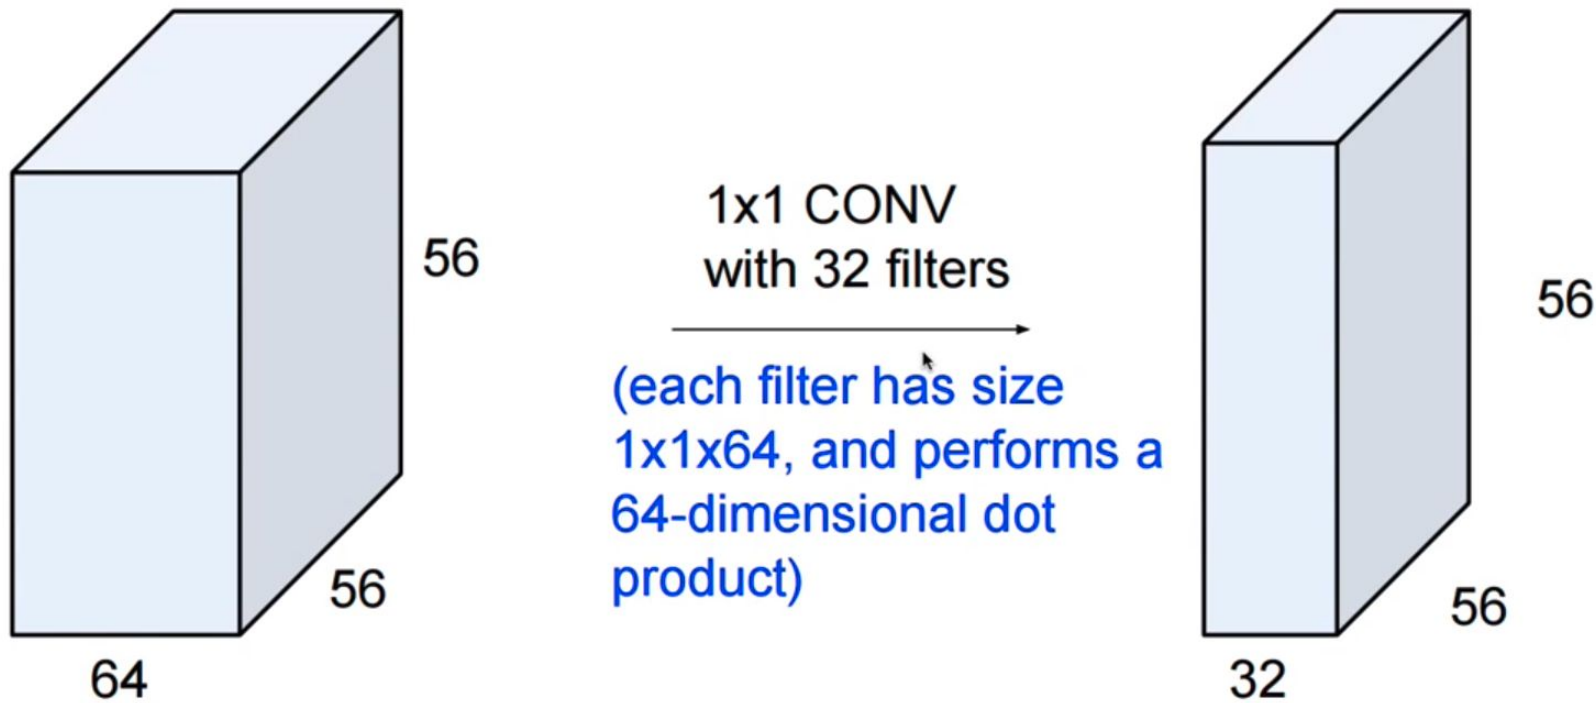
Usually inserted after Fully Connected or Convolutional layers, and before nonlinearity.

$$\hat{x}^{(k)} = \frac{x^{(k)} - \mathbb{E}[x^{(k)}]}{\sqrt{\text{Var}[x^{(k)}]}}$$

# Fully Connected Layer




# Convolution 1x1





# Regularization: Dropout

How can this possibly be a good idea?

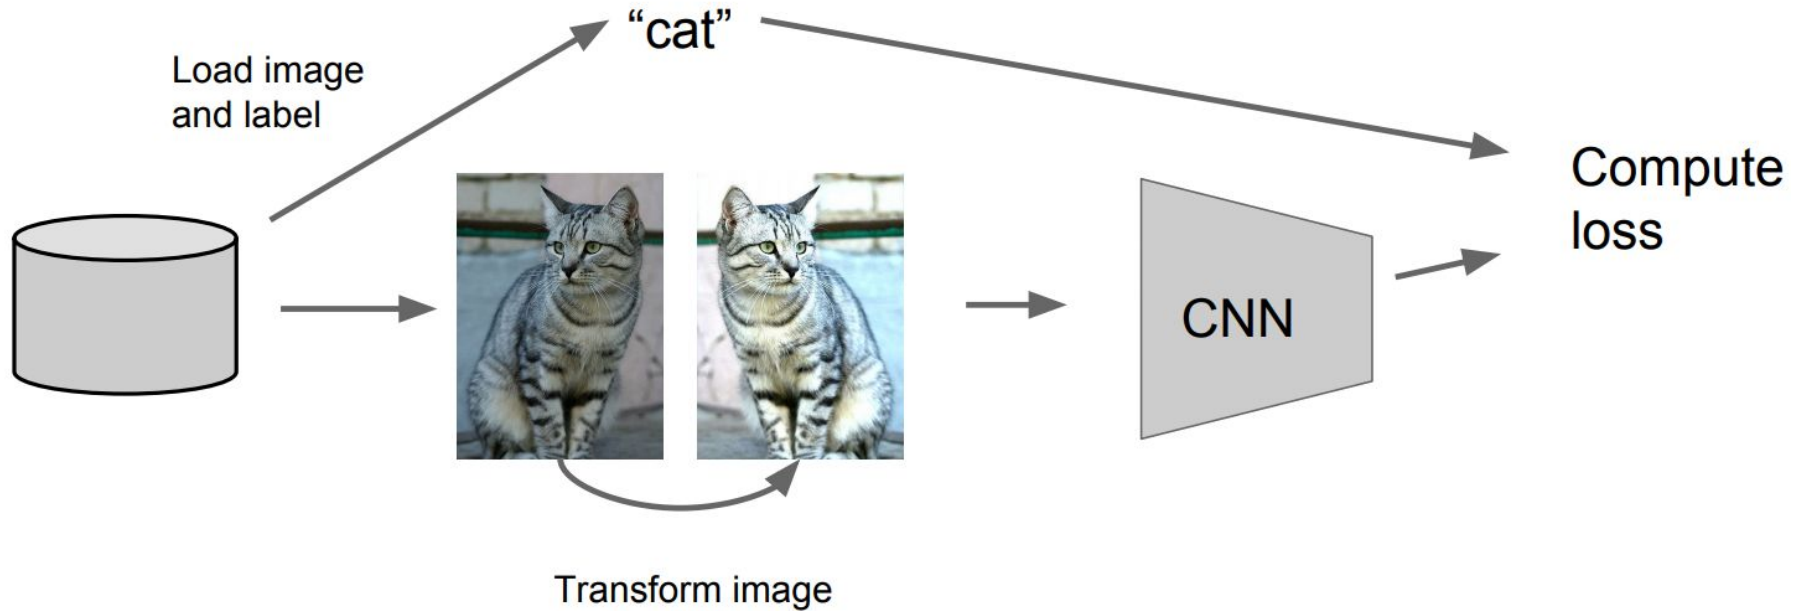


Forces the network to have a redundant representation;  
Prevents co-adaptation of features






# Regularization: Data Augmentation



# Activation Functions


## Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




## tanh

$$\tanh(x)$$




## ReLU

$$\max(0, x)$$



## Leaky ReLU

$$\max(0.1x, x)$$




## Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

## ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



SCORES

SOFTMAX

PROBABILITIES

$$y \begin{bmatrix} 2.0 \\ 1.0 \\ 0.1 \end{bmatrix}$$

$$S(y_i) = \frac{e^{y_i}}{\sum_j e^{y_j}}$$

$$\begin{aligned} &\rightarrow p = 0.7 \\ &\rightarrow p = 0.2 \\ &\rightarrow p = 0.1 \end{aligned}$$

# CROSS-ENTROPY

$S(Y)$

0.7
0.2
0.1

$L$

1.0
0.0
0.0

$$D(S, L) = - \sum_i L_i \log(S_i)$$


$$D(S, L) \neq D(L, S)$$

# Arquitecturas

- LeNet
- AlexNet
- ZFNet
- VGG
- GoogleNet
- ResNet
- InceptionResNet
- NASNet
- Xception
- MobileNet
- SeNet

# Review: LeNet-5

[LeCun et al., 1998]




Conv filters were 5x5, applied at stride 1

Subsampling (Pooling) layers were 2x2 applied at stride 2

i.e. architecture is [CONV-POOL-CONV-POOL-FC-FC]

# Modelos CNN Clasificación

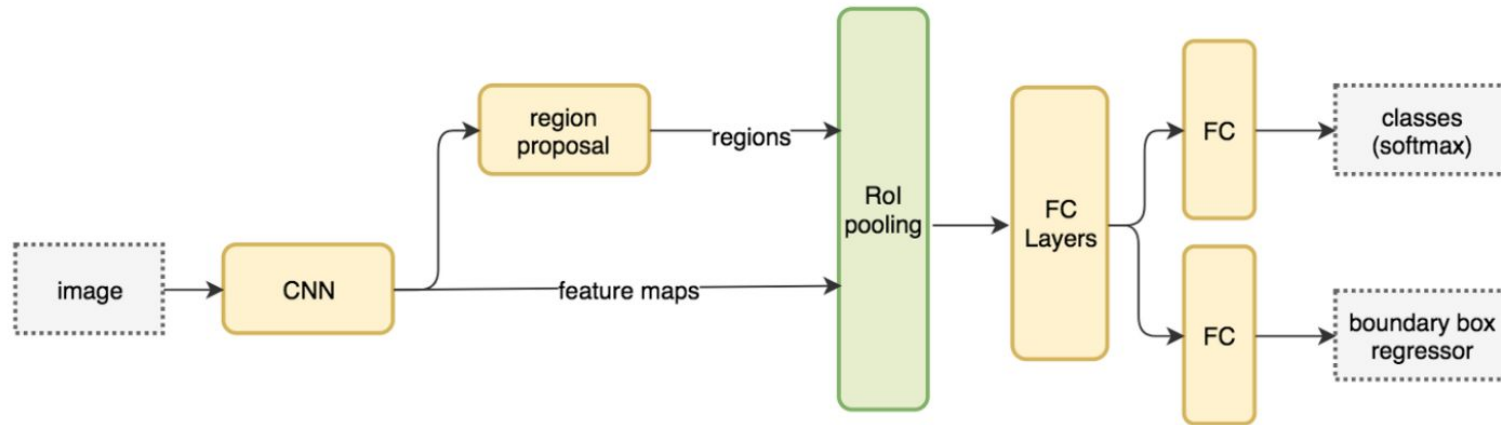


Human error (5.1%)  
surpassed in 2015


- **AlexNet (2012): First CNN (15.4%)**
  - 8 layers
  - 61 million parameters
- **ZFNet (2013): 15.4% to 11.2%**
  - 8 layers
  - More filters. Denser stride.
- **VGGNet (2014): 11.2% to 7.3%**
  - Beautifully uniform:  
3x3 conv, stride 1, pad 1, 2x2 max pool
  - 16 layers
  - 138 million parameters
- **GoogLeNet (2014): 11.2% to 6.7%**
  - Inception modules
  - 22 layers
  - 5 million parameters  
(throw away fully connected layers)
- **ResNet (2015): 6.7% to 3.57%**
  - More layers = better performance
  - 152 layers
- **CUImage (2016): 3.57% to 2.99%**
  - Ensemble of 6 models
- **SENet (2017): 2.99% to 2.251%**
  - Squeeze and excitation block: network is allowed to adaptively adjust the weighting of each feature map in the convolutional block.

# Object Detection / Localization

Region-Based Methods | Shown: Faster R-CNN



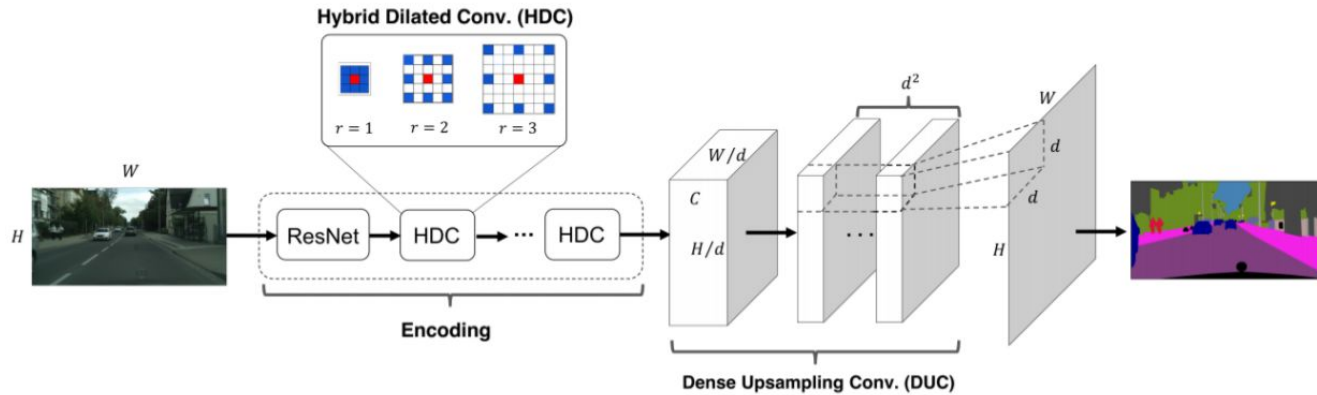
```
ROIs = region_proposal(image)
for ROI in ROIs
    patch = get_patch(image, ROI)
    results = detector(patch)
```







# Semantic Segmentation



# Transfer Learning

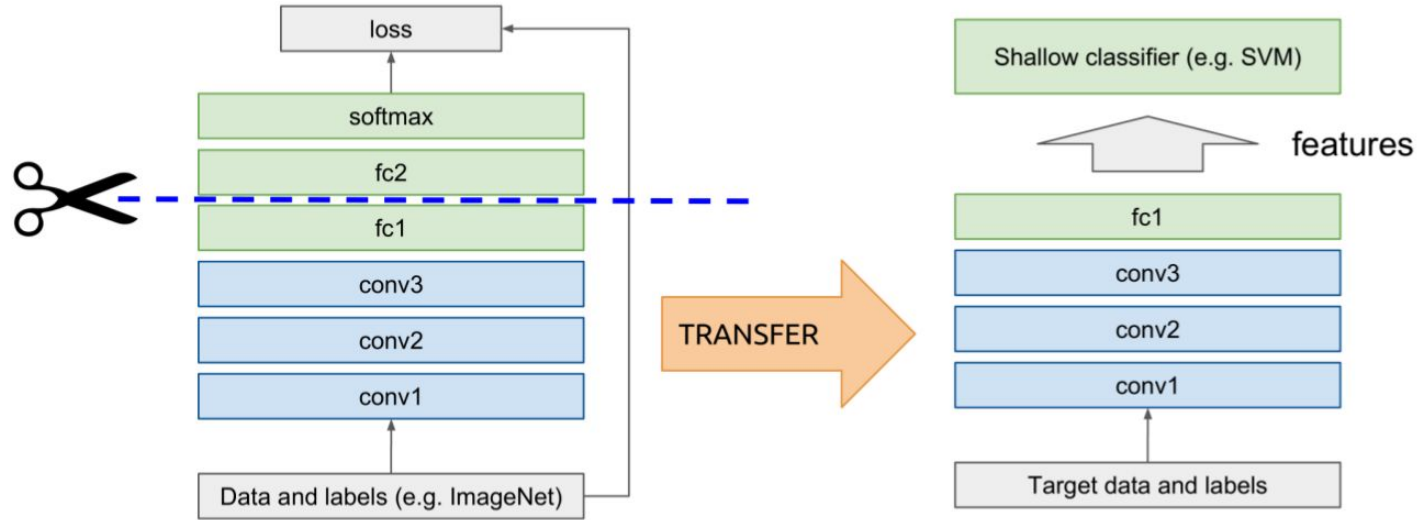
“You need a lot of a data if you want to  
train/use CNNs”

# Transfer Learning

“You need a lot of data if you want to train/use CNNs”



**BUSTED**

# Transfer Learning




- Fine-tune a pre-trained model
- Effective in many applications: computer vision, audio, speech, natural language processing

# Guide for How to Use Transfer Learning

Size of Data Set	LARGE	Fine-tune	Fine-tune or Retrain
	SMALL	End of ConvNet	Start of ConvNet
		<div>SIMILAR</div>	<div>DIFFERENT</div>
		Similarity to Training Data	

# Pre-trained Convolutional Neural Network



# Case: Small Data Set, Similar Data

Size of Data Set

LARGE

SMALL

Fine-tune

Fine-tune or  
Retrain

End of ConvNet

Start of ConvNet

SIMILAR



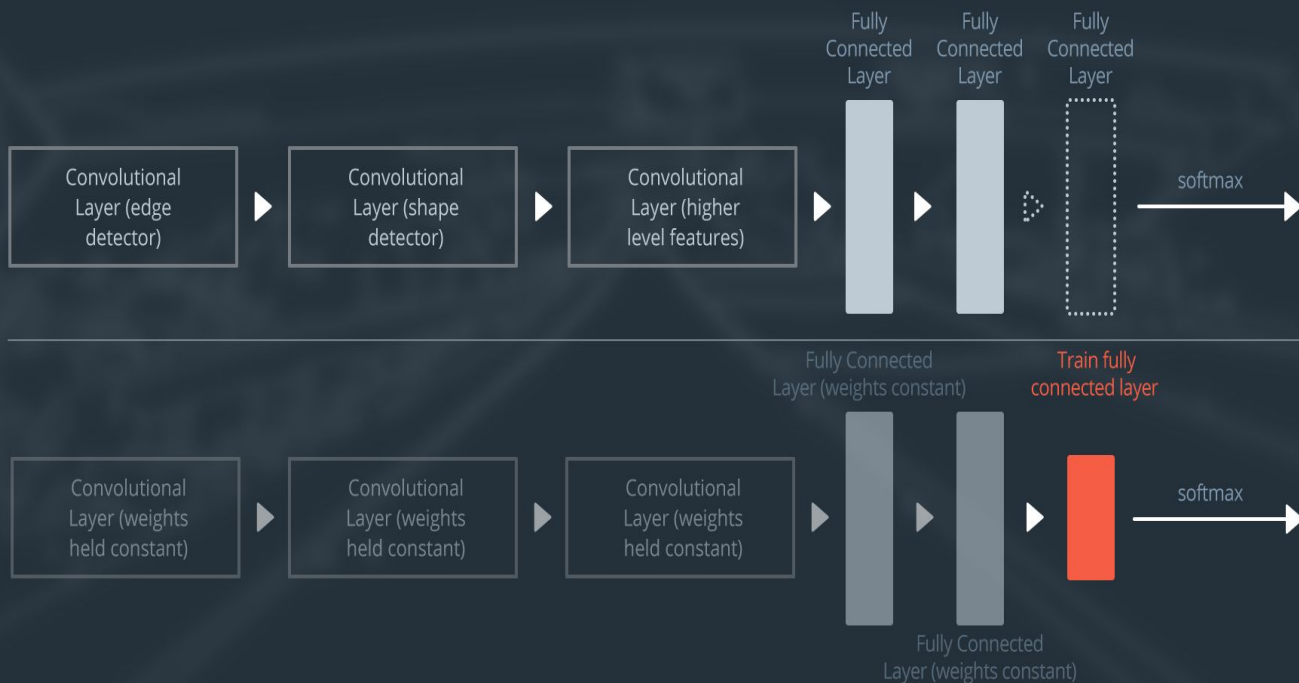
DIFFERENT



Similarity to Training Data







# Case: Small Data Set, Similar Data




- Quitar la parte final de la red, la última FC. Agregar un FC que matchee con la cantidad de clases e inicializar pesos random. Entrenar solo los pesos de la última FC.

- Congelar los pesos de las capas anteriores para evitar overfitting. Ya que son datasets similares, las imágenes de cada uno van a tener similares high level features. Además, las capas de la red pre-entrenada contiene información relevante y debe mantenerse.


# Case: Small Data Set, Different Data

Size of Data Set			
		LARGE	SMALL
	LARGE	Fine-tune	Fine-tune or Retrain
	SMALL	End of ConvNet	Start of ConvNet
		SIMILAR  	DIFFERENT  
		Similarity to Training Data	

# Case: Small Data Set, Different Data



Train fully connected layer





- Quitar la mayoría de capas pre-entrenadas hasta llegar a las características de bajo nivel . Agregar un FC que matchee con la cantidad de clases e inicializar pesos random. Entrenar solo los pesos de la última FC.

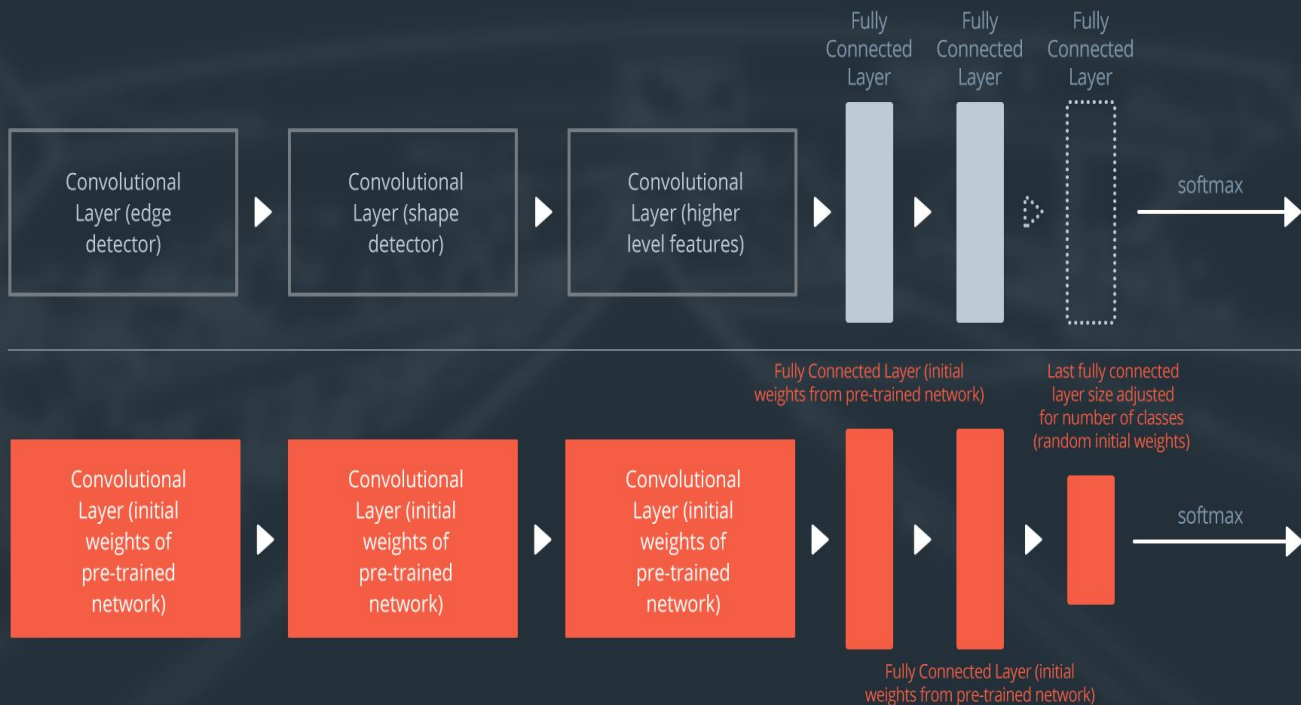
- Debido a que el dataset es pequeño existe la posibilidad de overfitting. Para ello se congelan los pesos.

- Al ser diferente data, no comparten características de alto nivel. Por lo que solo usaremos de bajo nivel.

# Case: Large Data Set, Similar Data

Size of Data Set	LARGE	Fine-tune	Fine-tune or Retrain
	SMALL	End of ConvNet	Start of ConvNet
		SIMILAR 	DIFFERENT 
		Similarity to Training Data	

# Case: Large Data Set, Similar Data





- Quitar la parte final de la red, la última FC. Agregar un FC que matchee con la cantidad de clases e inicializar pesos random. Entrenar solo los pesos de la última FC.

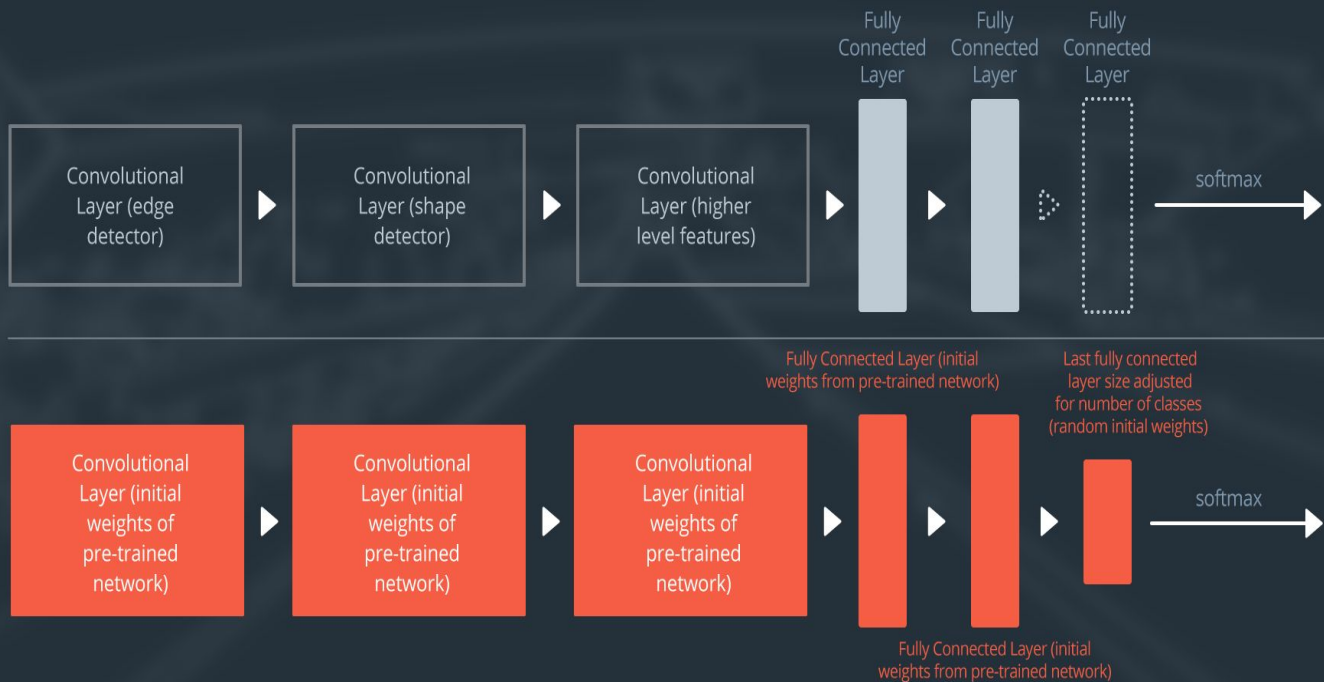
- Inicializa los pesos de las capas de red pre-entrenada con los pesos pre-entrenados sin congelarlos y re-entrenar el modelo.

- Debido a que los datasets son de clases similares y comparten los mismo high level features, se usa toda la red.

# Case: Large Data Set, Different Data

Size of Data Set	LARGE	Fine-tune	Fine-tune or Retrain
	SMALL	End of ConvNet	Start of ConvNet
		SIMILAR 	DIFFERENT 
		Similarity to Training Data	

# Case: Large Data Set, Different Data



- Quitar la parte final de la red, la última FC. Agregar un FC que matchee con la cantidad de clases e inicializar pesos random. Entrenar solo los pesos de la última FC.

- Inicializa los pesos de las capas de red pre-entrenada con pesos pre-entrenados como el caso anterior. En caso no fuera exitoso, iniciar los pesos en random y re-entrenar el modelo.

# Notebook

[https://github.com/leo2105/Keras\\_Introduction/blob/master/5\\_CustomClassifier.ipynb](https://github.com/leo2105/Keras_Introduction/blob/master/5_CustomClassifier.ipynb)

## Issues:

- Usar python 3.6, pip install python=3.6
- Instalar matplotlib, pip install matplotlib
- Crear file data2



## Referencias:

- <http://cs231n.stanford.edu/>
- <https://deeplearning.mit.edu/>
- <https://www.udacity.com/course/deep-learning--ud730>

Gracias!